

## 1. Basic Initialization

- 1.1. Load the expanded lab initial configs onto devices R1, R3, XR3, XR4, and R11 – R20, which contain basic IPv4 & IPv6 addressing per the diagram. Once complete, ensure that all devices can ping their directly connected neighbors with both IPv4 & IPv6.

## 2. Carrier Supporting Carrier (CSC)

- 2.1. R14 and R19 are the final CE routers that will be using the CSC network for transport. R14 is dual homed to its MPLS L3VPN provider via R12 and R13. R19's MPLS L3VPN provider is R17. Configure EIGRP AS 1419 on these devices, and advertise their IPv4 and IPv6 Loopback0 interfaces.
- 2.2. R11, R12, R13, R16, R17, and R18 are in BGP AS 1000. Use IS-IS Level 1 routing on all links and Loopback0 interfaces in AS 1000 to provide IGP transport for both IPv4 and IPv6 within the two sites.
- 2.3. Configure BGP as follows to provide CSC transport:
  - 2.3.1. R11 should peer IPv4 Unicast BGP with XR4
  - 2.3.2. R16 should peer IPv4 Unicast BGP with R1
  - 2.3.3. R12, R13, & R17 should peer full mesh VPNv4 and VPNv6 BGP
- 2.4. Redistribute wherever necessary. Once the CSC network is completed, you should have both IPv4 and IPv6 connectivity between R14 and R19, with R14 load balancing traffic between both R12 and R13.

## 3. L3VPN Extranet

- 3.1. R20 is hosting shared services for AS 123's L3VPN customers. Configure Extranet routing using BGP so that R6, R7, R8, & R10 are all able to reach the IPv4 and IPv6 Loopback0 addresses of R20. Traffic should not leak between R8/R10's site and R6/R7's site. Use only a single BGP peering on R20 to XR3 to accomplish this.

## 4. L3VPN Internet Access

- 4.1. R3 connects to the Internet via R15. Configure the network so that R6, R7, R8, and R10 can all ping the IPv4 Loopback0 address of R15, while following these restrictions:
  - 4.1.1. Do not modify any routing configuration on R15.
  - 4.1.2. R3's link to R15 should remain in the global routing table.
  - 4.1.3. Traffic should not leak between R8/R10's site and R6/R7's site.

## 5. Unified MPLS

5.1. Remove the redistribution of OSPFv2 and IPv4 IS-IS on both R2 and XR2.

5.2. Modify your BGP peerings in AS 123 as follows:

5.2.1. R2 and XR2 should be route reflectors for IPv4 unicast to R1, R3, XR1, and XR3.

5.2.2. Advertise the IPv4 Loopback0 networks of all six of these routers into IPv4 Unicast BGP.

5.2.3. Once complete, any of the L3VPN traffic transiting over the core should still be successful, e.g. R6 to R7 IPv4 and IPv6 transport.

## 6. LISP

6.1. Modify the Customer 1 and Customer 2 sites PE-CE routing as follows:

6.1.1. Remove the VRF aware OSPFv2 process for Customer 1 & 2 on R1, R3, XR1, and XR3.

6.1.2. Configure R4, R5, and R7 with static default routes to their attached PE routers. Advertise these default routes into OSPFv2 on R4 and R5. Configure R20 with a static default route to XR3.

6.1.3. Advertise the following links into the VRF aware BGP process:

6.1.3.1. R1's link to R5

6.1.3.2. R3's link to R7

6.1.3.3. XR1's link to R4

6.1.3.4. XR3's link to R7

6.1.3.5. XR3's link to R20

6.1.4. Ensure that you can ping IPv4 from R4, R5, & R7 to R20.

6.2. Configure LISP on R4, R5, R7 & R20 as follows:

6.2.1. R20 will be the Map Server and Map Resolver (MS/MR)

6.2.2. R4, R5, & R7 will be the Ingress and Egress Tunnel Routers (xTRs)

6.2.3. R4 and R5 should register their links to the PE's as the Routing Locators (RLOCs) to R20 as "SITE1".

6.2.4. R7 should register its links to both PEs as RLOCs to R20 as "SITE2".

6.2.5. Once complete, R6 and R7 should have IP reachability to each other's Loopback0 networks.

6.3. Modify the LISP configuration so that R4 accepts twice as much traffic for Site 1 as R5.

## 7. MPLS over GRE

- 7.1. Create a new Loopback1 interface on R5 and R7 with addresses 1.5.5.5/32 and 1.7.7.7/32. Configure a GRE tunnel between them sourced and unnumbered from Loopback0. Enable a new EIGRP process with AS 57 on the tunnel and the new Loopbacks. Enable LDP on the tunnel. Once complete you should have an LSP between the new Loopbacks.

## 8. L2VPN

- 8.1. Configure new IPv4 addresses 192.168.0.4/24 on R4's layer 2 link to R7, and address 192.168.0.8/24 on R8's layer 2 link to R5. Configure a point-to-point L2VPN circuit between R5 and R7 to transport the traffic between R4 and R8. Once complete, R4 and R8 should be able to ping each other over the layer 2 circuit.
- 8.2. Configure new IPv4 addresses 172.16.0.2/24, 172.16.0.9/24, and 172.16.0.22/24 on R2, R9, and XR2's additional links per the L2VPN diagram.
- 8.3. Configure BGP L2VPN peerings as follows:
  - 8.3.1. R1 to R3
  - 8.3.2. R2 to R3
  - 8.3.3. R3 to XR3
  - 8.3.4. XR1 to XR3
- 8.4. Establish a full mesh of pseudowires between PEs R1, R2, and XR1. Use BGP for signaling the VPLS membership.

## 9. MPLS Traffic Engineering

- 9.1. Enable OSPFv2 area 0 on all links between R1, R2, XR1, and XR2.
- 9.2. Configure an MPLS TE tunnel from R3 to XR1's Loopback using any available dynamic path in the core that has 25Mbps of bandwidth available. Dynamically route XR1's Loopback via the tunnel.  
fs
- 9.3. Configure an MPLS TE tunnel from XR3 to XR2's Loopback so that it follows the path of XR3 > R3 > R2 > XR2. If this path is unavailable, fall back to any other available path. Use static routing out the tunnel to reach **R1**. Once complete you should be able to ping from XR3's interface in the Customer Site 2 VRF to R1's interface in the Customer Site 1 VRF.
- 9.4. Modify R3's tunnel to XR1 so that it routes from R3 > R2 > XR2 > XR1. Do not use an explicit path for this or change any IGP metrics.
- 9.5. Configure R2 to protect R3's tunnel to XR1 in the case that R2's link to XR2 goes down. If this failure occurs, it should be detected within 500ms and traffic should re-route from R2 to R1 to XR2.

## 10.fsMulticast

- 10.1. Revert your configs to the VIRL file "csc.working.virl"
- 10.2. Enable multicast routing in VRF A on R2, R8, R10 & XR1. Use R2 RP for this customer. Use the BGP MDT AFI to advertise that R2 and XR1 are participating in multicast routing for customer A. Enable PIM SSM in the core where necessary to transport multicast traffic. Join the multicast group 228.8.8.8 on R8's Loopback0. Once complete, R10 should be able to ping this address when sourcing traffic from its own Loopback0 interface.
- 10.3. Enable multicast routing in VRF B on R1, R3, R4, R5, R6, R7, XR1, & XR3. Use R7's Loopback0 as the BSR and Candidate RP. Configure MLDP with R2 as the root of the tree and XR2 as the backup root of the tree for customer VRF B. Join the multicast group 226.6.6.6 on R6's Loopback0. Once complete, R7 should be able to ping this address when sourcing traffic from its own Loopback0 interface.